

Oct 06, 11 14:15

cu_nbody.cu

Page 1/8

```

/*
 * This is a commented version of the original N-Body
 * program, some of the internal variable has a different
 * but better readable name. The structure and principle
 * of this program has no big change.
 *
 * Commented By Daniel Zhang <danielzhang0212@gmail.com>
 */

#include <stdio.h>
#include <cmath>
#include <cassert>
#include <cutil.h>
#include <omp.h>
#include "cuda_pointer.h"

/*
 * define some constants, like
 * the number of thread of a block,
 * j blocks, and i blocks
 */
#define NTHREAD 64 // 64, 96, 128 or 192
#define NJBLOCK 16 // 8800GTS/512 has 16
#define NIBLOCK 16 // 16 or 32
#define NIMAX (NTHREAD * NIBLOCK) // 1024
#define NBMAX 128 // NNB per block

/*
 * user define vector
 */
template <class T>
struct myvector{
    int num;
    T *val;
    myvector(){
        num = 0;
        val = NULL;
    }
    ~myvector(){
        delete [] val;
    }
    void clear(){
        num = 0;
    }
    void reserve(size_t count){
        val = new T[count];
    }
    void free(){
        delete [] val;
    }
    void push_back(const T &t){
        val[num++] = t;
    }
    size_t size(){
        return num;
    }
    T &operator[](int i){
        return val[i];
    }
};

/*
 * define the time related functions
 * in order to record the calculation time.
 */
#include <sys/time.h>

```

Thursday October 06, 2011

cu_nbody.cu

Oct 06, 11 14:15

cu_nbody.cu

Page 2/8

```

static double get_wtime(){
    struct timeval tv;
    gettimeofday(&tv, NULL);
    return tv.tv_sec + 1.e-6 * tv.tv_usec;
}

static double time_send, time_grav;
static long long numInter

/*
 * define the structure of particle J
 * contain: position, velocity, mass and pad
 */
struct Jparticle{
    float3 pos;           // position
    float3 vel;           // velocity
    float mass;           // mass
    float pad;            // TODO:don't know!
    Jparticle() {}
    Jparticle(double mj, double posj[3], double velj[3]){
        pos.x = posj[0];
        pos.y = posj[1];
        pos.z = posj[2];
        mass = mj;
        vel.x = velj[0];
        vel.y = velj[1];
        vel.z = velj[2];
    }
};

/*
 * define the structure of particle I
 * contain: position, velocity, pad, and h2
 */
struct Iparticle{
    float3 pos;           // position
    float3 vel;           // velocity
    float pad;            // TODO:don't know
    float h2;             // Threshold to dicide a neighbor
    Iparticle() {}
    Iparticle(double h2i, double posi[3], double veli[3]){
        pos.x = posi[0];
        pos.y = posi[1];
        pos.z = posi[2];
        h2 = h2i;
        vel.x = veli[0];
        vel.y = veli[1];
        vel.z = veli[2];
    }
};

/*
 * define the structure of the force
 * between two particles, including jerk
 * force, acceleration, potential, and neighborhood
 */
struct Force{
    float3 acc;           // acceleration
    float3 jrk;           // jerk force
    float pot;            // potential
    int nnb;              // number of neighbor particles ONE BYTE
    unsigned short neib[NBMAX]; // 24 words
    __device__ Force(){ // created inside a kernel
        acc.x = acc.y = acc.z = 0.f;
        jrk.x = jrk.y = jrk.z = 0.f;
        pot = 0.f;
        nnb = 0;
    }
};

```

1/4

Oct 06, 11 14:15

cu_nbody.cu

Page 3/8

```

    }
};

/*
 * calculate the FORCE contents
 * between two particles.
 * using the position and velocity of
 * two particles.
 */
__device__ void force_kernel(
    const int id,
    const Iparticle &ip,
    const Jparticle &jp,
    Force &force)
{
    float dx = jp.pos.x - ip.pos.x;    // position diff of x
    float dy = jp.pos.y - ip.pos.y;    // position diff of y
    float dz = jp.pos.z - ip.pos.z;    // position diff of z
    float dvx = jp.vel.x - ip.vel.x;   // velocity diff of x
    float dvy = jp.vel.y - ip.vel.y;   // velocity diff of y
    float dvz = jp.vel.z - ip.vel.z;   // velocity diff of z

    float dist_square = dx*dx + dy*dy + dz*dz;

    float inverse = rsqrtf(dist_square); // used as denominator

    if(dist_square < ip.h2) // if the distance is smaller than a value
    {
        // add this particle into neighbour list
        force.neib[force.nnb & (NBMAX-1)] = (unsigned)j; // do mo
        // NBMAX-1=(11111111) in binary */
        force.nnb++;

        /* neighbor's force does not count */
        inverse = 0.f;
    }

    float inverse_square = inverse * inverse;
    float potential = jp.mass * inverse; // potential
    float acc_para = mass_inverse * inverse; // accelerate para

    rv = -3.f * inverse_square * (dx*dvx + dy*dvy + dz*dvz);

    /* the final calculation of the FORCE */
#ifdef POTENTIAL
    force.pot += potential;
#endif
    force.acc.x += acc_para * dx;
    force.acc.y += acc_para * dy;
    force.acc.z += acc_para * dz;
    force.jrk.x += acc_para * (dvx + rv * dx);
    force.jrk.y += acc_para * (dvy + rv * dy);
    force.jrk.z += acc_para * (dvz + rv * dz);

    // FUNCTION DONE
}

/*
 * Main kernel of the n-body forces calculation
 */
__global__ void KERNEL(
    int nbody, // number of bodies
    Iparticle ips[],
    Jparticle jps[],
    Force forces[][NJBLOCK])
{
    int i_bid = blockIdx.x;

```

Oct 06, 11 14:15

cu_nbody.cu

Page 4/8

```

    int j_bid = blockIdx.y;
    int tid = threadIdx.x;
    int i_addr = tid + NTHREAD * i_bid;
    int j_start=(nbody/NJBLOCK) * j_bid;
    int j_end=(nbody/NJBLOCK) * j_bid +1 ;

    Iparticle ip = ips[i_addr]; // fetch the i particle
    Force fo;

    /*
     * loop every block of j particle,
     * where variable j stand for the first
     * element of the Jblock
     */
    for(int j=j_start;j<j_end;j+=NTHREAD)
    {
        __shared__ Jparticle jpshare[NTHREAD]; // use shared memory
        __syncthreads();

        /* copy data into shared memory */
        float4 *src = (float4 *) &jps[j];
        float4 *dst = (float4 *) jpshare;
        dst[tid] = src[tid];
        dst[NTHREAD+tid] = src[NTHREAD+tid];
        /* TODO why copy 2 blocks?
        */
        __syncthreads();

        /*
         * calculate the I particle's force,
         * interactive with every J particle
         * in the whole system.
         */
        if(j_end-j<NTHREAD)
        {
            for(int partId=0;partId<jend-j;partId++)
            {
                Jparticle jp = jpshare[partId];
                force_kernel(partId, ip, jp, fo);
            }
        }
        else
        {
#pragma unroll
            for(int partId=0;partId<NTHREAD;partId++)
            {
                Jparticle jp = jpshare[partId];
                force_kernel(partId, ip, jp, fo);
            }
        }
    }
    /*
     * the fo's calculation is complete,
     * save it into buffer. fo stand for
     * the i's Iparticle interactive with the
     * j_bid's block of Jparticle
     */
    forces[i_addr][j_bid] = fo;
}

/*
 * The functions below are CPU functions
 */
static cudaPointer <Jparticle> jps;
static cudaPointer <Iparticle> ips;

```

Oct 06, 11 14:15	cu_nbody.cu	Page 5/8
------------------	--------------------	----------

```

static cudaPointer <Force[NJBLOCK]> forces;

#define MAX_CPU 1
static myvector<int> nblast[MAX_CPU];
static int nbody, nbodymax;

/*
 * the initial and end functions
 */
void GPUNB_open(int nbmax){
    time_send = time_grav = 0.0;
    numInter = 0;

    jpbuff.allocate(nbmax + NTHREAD);
    ipbuff.allocate(NIMAX);
    fobuff.allocate(NIMAX);
    nbodymax = nbmax;
#pragma omp parallel
    {
        int tid = 0;
        nblast[tid].reserve(nbmax);
    }
}
void GPUNB_close(){
    jpbuff.free();
    ipbuff.free();
    fobuff.free();
    nbodymax = 0;
}

#ifdef PROFILE
    fprintf(stderr, "*****\n");
    fprintf(stderr, "time send : %f sec\n", time_send);
    fprintf(stderr, "time grav : %f sec\n", time_grav);
    fprintf(stderr, "%f Gflops (gravity part only)\n", 60.e-9 * numInter / t
ime_grav);
    fprintf(stderr, "*****\n");
#endif
}
void GPUNB_send(
    int nj,
    double mj[],
    double xj[][3],
    double vj[][3]){
    time_send -= get_wtime();
    nbody = nj;
    assert(nbody <= nbodymax);
    for(int j=0; j<nj; j++){
        jpbuff[j] = Jparticle(mj[j], xj[j], vj[j]);
    }

    jpbuff.htod(nj);
    time_send += get_wtime();
}

/*
 * MAIN FUNCTION of this program.
 * Call the kernel.
 */
void GPUNB_regf(
    int nI, // number of I particle
    double h2[], // one para of Iparticle
    double xi[][3], // position vector
    double vi[][3], // velocity vector
    double acc[][3], // acceleration vecotor
    double jrk[][3], // force vector
    double pot[], // potential
    int lmax, // TODO

```

Oct 06, 11 14:15	cu_nbody.cu	Page 6/8
------------------	--------------------	----------

```

    int nbMax, // max number of neighbo
    int *listbase) // TODO
{
    /* get initial time */
    time_grav -= get_wtime();
    numInter += ni * nbody;
    assert(0 < ni && ni <= NIMAX)

    for(int i=0;i<nI; i++)
    {
        ips[i] = Iparticle(h2[i],xi[i],vi[i]);
    }

    /* load the particles into device memory */
    ipbuff.htod(ni);

    /* start the kernel */
    int niblock = 1 + (nI - 1) / NTHREAD;
    dim3 grid(niblock, NJBLOCK, 1);
    dim3 threads(NTHREAD, 1,1);
    KERNEL <<< grid, threads >>> (nbody,ips,jps,forces)

    forces.dtoh(nI);
    // reduction phase

#pragma omp parallel for
    for(int i=0;i<nI;i++)
    {
        // TODO I think "tid" should be "i"
        // For it has never changed in
        // every loop.
        int tid=0;
        double ax=0,ay=0,az=0; // acceleration
        double jx=0,jy=0,jz=0; // jerk force

#ifdef POTENTIAL
        double poti=0;
#endif
        for(int jblock=0;jblock<NJBLOCK;jblock++)
        {
            Force &fo = forces[i][jblock];
            ax += fo.acc.x;
            ay += fo.acc.y;
            az += fo.acc.z;
            jx += fo.jrk.x;
            jy += fo.jrk.y;
            jz += fo.jrk.z;

#ifdef POTENTIAL
            poti += fo.pot;
#endif
        }
        /* save the sum of individual result */
        acc[i][0] = ax;
        acc[i][1] = ay;
        acc[i][2] = az;
        jrk[i][0] = jx;
        jrk[i][1] = jy;
        jrk[i][2] = jz;

#ifdef POTENTIAL
        pot[i] = poti;
#endif

        /* TODO something about neighbors */
        bool overflow = false;
        for(int jblock=0;jblock<NJBLOCK;jblock++)
        {
            Force &fo = forces[i][jblock];
            int jstart = (nbody / NJBLOCK) * jblock;

```

Oct 06, 11 14:15

cu_nbody.cu

Page 7/8

```

        if(fo.nnb <= NBMAX)
        {
            for(int k=0;k<f.nnb;k++)
            {
                int nb = fo.neib[k];
                while(nb<jstart)
                    nb += (1<<16);
                nblist[tid].push_back(nb);
            }
        }
        else
        {
            overflow = true;
        }
    }

    int * nnbp = listbase + lmax * i;          // number of neighbor pa
rticles
    int * nblistp = nnbp +1;
    int nnb = nblist[tid].size();
    if(nnb > nbmax)
        overflow =true;

    if(overflow)
    {
        *nnbp = -1;
    }
    else
    {
        *nnbp = nnb;
        for(int k=0;k<nnb;k++)
        {
            nblistp[k]=nblist[tid][k];
        }
    }

}

}

/* TODO combine programming between C and C++ */
extern "C" {
    void gpunb_open_(int *nbmax){
        GPUNB_open(*nbmax);
    }
    void gpunb_close_(){
        GPUNB_close();
    }
    void gpunb_send_(
        int *nj,
        double mj[],
        double xj[][3],
        double vj[][3]){
        GPUNB_send(*nj, mj, xj, vj);
    }
    void gpunb_regf_(
        int *ni,
        double h2[],
        double xi[][3],
        double vi[][3],
        double acc[][3],
        double jrk[][3],
        double pot[],
        int *lmax,
        int *nbmax,
        int *list){ // list[][lmax]
        GPUNB_regf(*ni, h2, xi, vi, acc, jrk, pot, *lmax, *nbmax, list);
    }
}

```

Oct 06, 11 14:15

cu_nbody.cu

Page 8/8

```

    }
}

```